

# assignment01\_MuleyTushar

December 2, 2021

Name: Muley, Tushar Assignment: Week1-Assignment 1.1 Examples mnist\_mlp and Pyspark

Example: mnist\_mlp

```
[1]: '''Trains a simple deep NN on the MNIST dataset.  
  
Gets to 98.40% test accuracy after 20 epochs  
(there is *a lot* of margin for parameter tuning).  
2 seconds per epoch on a K520 GPU.  
'''  
  
from tensorflow import keras  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout  
from tensorflow.keras.optimizers import RMSprop  
  
batch_size = 128  
num_classes = 10  
epochs = 20  
  
# the data, split between train and test sets  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
  
x_train = x_train.reshape(60000, 784)  
x_test = x_test.reshape(10000, 784)  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
x_train /= 255  
x_test /= 255  
print(x_train.shape[0], 'train samples')  
print(x_test.shape[0], 'test samples')  
  
# convert class vectors to binary class matrices  
y_train = keras.utils.to_categorical(y_train, num_classes)  
y_test = keras.utils.to_categorical(y_test, num_classes)  
  
model = Sequential()
```

```

model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11493376/11490434 [=====] - 1s 0us/step

60000 train samples

10000 test samples

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Epoch 1/20

469/469 [=====] - 5s 11ms/step - loss: 0.2454 - accuracy: 0.9234 - val\_loss: 0.1155 - val\_accuracy: 0.9644

Epoch 2/20  
469/469 [=====] - 5s 10ms/step - loss: 0.1030 - accuracy: 0.9685 - val\_loss: 0.0830 - val\_accuracy: 0.9752

Epoch 3/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0737 - accuracy: 0.9780 - val\_loss: 0.0793 - val\_accuracy: 0.9774

Epoch 4/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0593 - accuracy: 0.9822 - val\_loss: 0.0803 - val\_accuracy: 0.9784

Epoch 5/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0493 - accuracy: 0.9854 - val\_loss: 0.0757 - val\_accuracy: 0.9801

Epoch 6/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0448 - accuracy: 0.9875 - val\_loss: 0.0801 - val\_accuracy: 0.9805

Epoch 7/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0376 - accuracy: 0.9890 - val\_loss: 0.0944 - val\_accuracy: 0.9812

Epoch 8/20  
469/469 [=====] - 4s 10ms/step - loss: 0.0332 - accuracy: 0.9898 - val\_loss: 0.0959 - val\_accuracy: 0.9811

Epoch 9/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0321 - accuracy: 0.9906 - val\_loss: 0.0869 - val\_accuracy: 0.9827

Epoch 10/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0292 - accuracy: 0.9916 - val\_loss: 0.0924 - val\_accuracy: 0.9826

Epoch 11/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0279 - accuracy: 0.9922 - val\_loss: 0.0899 - val\_accuracy: 0.9827

Epoch 12/20  
469/469 [=====] - 4s 10ms/step - loss: 0.0261 - accuracy: 0.9928 - val\_loss: 0.0996 - val\_accuracy: 0.9811

Epoch 13/20  
469/469 [=====] - 4s 9ms/step - loss: 0.0256 - accuracy: 0.9934 - val\_loss: 0.1012 - val\_accuracy: 0.9830

Epoch 14/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0225 - accuracy: 0.9940 - val\_loss: 0.1180 - val\_accuracy: 0.9815

Epoch 15/20  
469/469 [=====] - 4s 10ms/step - loss: 0.0228 - accuracy: 0.9936 - val\_loss: 0.1134 - val\_accuracy: 0.9828

Epoch 16/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0227 - accuracy: 0.9942 - val\_loss: 0.1130 - val\_accuracy: 0.9827

Epoch 17/20  
469/469 [=====] - 4s 9ms/step - loss: 0.0199 - accuracy: 0.9945 - val\_loss: 0.1196 - val\_accuracy: 0.9844

```

Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0190 -
accuracy: 0.9946 - val_loss: 0.1136 - val_accuracy: 0.9837
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0203 -
accuracy: 0.9947 - val_loss: 0.1448 - val_accuracy: 0.9824
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0184 -
accuracy: 0.9956 - val_loss: 0.1211 - val_accuracy: 0.9838
Test loss: 0.12112750113010406
Test accuracy: 0.9837999939918518

```

Example Pyspark

```

[3]: #
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#

import sys
from random import random
from operator import add

from pyspark.sql import SparkSession

if __name__ == "__main__":
    """
        Usage: pi [partitions]
    """
    spark = SparkSession\
        .builder\
        .appName("PythonPi")\
        .getOrCreate()

    #partitions = int(sys.argv[1]) if len(sys.argv) > 1 else 2

```

```
partitions = 2
n = 100000 * partitions

def f(_):
    x = random() * 2 - 1
    y = random() * 2 - 1
    return 1 if x ** 2 + y ** 2 <= 1 else 0

count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).
↪reduce(add)
print("Pi is roughly %f" % (4.0 * count / n))

spark.stop()
```

Pi is roughly 3.142900

[ ]: