

# assignment05\_Muley\_Tushar\_Week3

December 16, 2021

**0.0.1 Name: Muley, Tushar**

**Assignment: Assignment 5 Week 3**

**Date: Dec 19, 2021**

## 0.1 Assignment 5.1

Implement the movie review classifier found in section 3.4 of Deep Learning with Python.

```
[1]: from keras.datasets import imdb
```

```
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>

17465344/17464789 [=====] - 1s 0us/step

<\_\_array\_function\_\_ internals>:5: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

C:\Users\Tushar\AppData\Roaming\Python\Python38\site-packages\tensorflow\python\keras\datasets\imdb.py:159:

VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

```
    x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])
```

C:\Users\Tushar\AppData\Roaming\Python\Python38\site-packages\tensorflow\python\keras\datasets\imdb.py:160:

VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

```
    x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

```
[2]: train_data[0]
```

[2] : [1,  
14,  
22,  
16,  
43,  
530,  
973,  
1622,  
1385,  
65,  
458,  
4468,  
66,  
3941,  
4,  
173,  
36,  
256,  
5,  
25,  
100,  
43,  
838,  
112,  
50,  
670,  
2,  
9,  
35,  
480,  
284,  
5,  
150,  
4,  
172,  
112,  
167,  
2,  
336,  
385,  
39,  
4,  
172,  
4536,  
1111,  
17,  
546,

38,  
13,  
447,  
4,  
192,  
50,  
16,  
6,  
147,  
2025,  
19,  
14,  
22,  
4,  
1920,  
4613,  
469,  
4,  
22,  
71,  
87,  
12,  
16,  
43,  
530,  
38,  
76,  
15,  
13,  
1247,  
4,  
22,  
17,  
515,  
17,  
12,  
16,  
626,  
18,  
2,  
5,  
62,  
386,  
12,  
8,  
316,  
8,

106,  
5,  
4,  
2223,  
5244,  
16,  
480,  
66,  
3785,  
33,  
4,  
130,  
12,  
16,  
38,  
619,  
5,  
25,  
124,  
51,  
36,  
135,  
48,  
25,  
1415,  
33,  
6,  
22,  
12,  
215,  
28,  
77,  
52,  
5,  
14,  
407,  
16,  
82,  
2,  
8,  
4,  
107,  
117,  
5952,  
15,  
256,  
4,

2,  
7,  
3766,  
5,  
723,  
36,  
71,  
43,  
530,  
476,  
26,  
400,  
317,  
46,  
7,  
4,  
2,  
1029,  
13,  
104,  
88,  
4,  
381,  
15,  
297,  
98,  
32,  
2071,  
56,  
26,  
141,  
6,  
194,  
7486,  
18,  
4,  
226,  
22,  
21,  
134,  
476,  
26,  
480,  
5,  
144,  
30,  
5535,

```
18,  
51,  
36,  
28,  
224,  
92,  
25,  
104,  
4,  
226,  
65,  
16,  
38,  
1334,  
88,  
12,  
16,  
283,  
5,  
16,  
4472,  
113,  
103,  
32,  
15,  
16,  
5345,  
19,  
178,  
32]
```

```
[3]: train_labels[0]
```

```
[3]: 1
```

```
[4]: max([max(sequence) for sequence in train_data])
```

```
[4]: 9999
```

**Decode one of these reviews back to English**

```
[5]: word_index = imdb.get_word_index()  
reverse_word_index = dict(  
    [(value, key) for (key, value) in word_index.items()])  
decoded_review = ' '.join(  
    [reverse_word_index.get(i - 3, '?') for i in train_data[0]])
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\\_word\\_index.json](https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json)

1646592/1641221 [=====] - 0s 0us/step

### Encoding the integer sequences into a binary matrix

```
[6]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
[7]: x_train[0]
```

```
[7]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
[8]: y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

### Building your network

```
[9]: from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential ([
    layers.Dense(16, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])
```

### Compile the model

```
[10]: model.compile(optimizer='rmsprop',
                    loss='binary_crossentropy',
                    metrics=['accuracy'])
```

### Validation

```
[11]: x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

### Training set

```
[12]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
```

```
validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 7s 136ms/step - loss: 0.5928 - accuracy: 0.7062 - val\_loss: 0.4317 - val\_accuracy: 0.8234

Epoch 2/20

30/30 [=====] - 1s 18ms/step - loss: 0.3378 - accuracy: 0.8964 - val\_loss: 0.3248 - val\_accuracy: 0.8748

Epoch 3/20

30/30 [=====] - 1s 18ms/step - loss: 0.2406 - accuracy: 0.9269 - val\_loss: 0.2798 - val\_accuracy: 0.8903

Epoch 4/20

30/30 [=====] - 1s 18ms/step - loss: 0.1856 - accuracy: 0.9417 - val\_loss: 0.2827 - val\_accuracy: 0.8864

Epoch 5/20

30/30 [=====] - 1s 26ms/step - loss: 0.1464 - accuracy: 0.9542 - val\_loss: 0.2961 - val\_accuracy: 0.8831

Epoch 6/20

30/30 [=====] - 1s 36ms/step - loss: 0.1197 - accuracy: 0.9647 - val\_loss: 0.2943 - val\_accuracy: 0.8870

Epoch 7/20

30/30 [=====] - 2s 43ms/step - loss: 0.1016 - accuracy: 0.9707 - val\_loss: 0.3267 - val\_accuracy: 0.8767

Epoch 8/20

30/30 [=====] - 1s 30ms/step - loss: 0.0840 - accuracy: 0.9766 - val\_loss: 0.3249 - val\_accuracy: 0.8811

Epoch 9/20

30/30 [=====] - 1s 23ms/step - loss: 0.0720 - accuracy: 0.9804 - val\_loss: 0.3547 - val\_accuracy: 0.8784

Epoch 10/20

30/30 [=====] - 1s 22ms/step - loss: 0.0586 - accuracy: 0.9853 - val\_loss: 0.3711 - val\_accuracy: 0.8799

Epoch 11/20

30/30 [=====] - 1s 18ms/step - loss: 0.0465 - accuracy: 0.9903 - val\_loss: 0.4031 - val\_accuracy: 0.8761

Epoch 12/20

30/30 [=====] - 1s 20ms/step - loss: 0.0369 - accuracy: 0.9924 - val\_loss: 0.4278 - val\_accuracy: 0.8746

Epoch 13/20

30/30 [=====] - 1s 21ms/step - loss: 0.0288 - accuracy: 0.9953 - val\_loss: 0.4565 - val\_accuracy: 0.8735

Epoch 14/20

30/30 [=====] - 1s 19ms/step - loss: 0.0266 - accuracy: 0.9946 - val\_loss: 0.4946 - val\_accuracy: 0.8702

Epoch 15/20

30/30 [=====] - 1s 19ms/step - loss: 0.0198 - accuracy: 0.9976 - val\_loss: 0.5417 - val\_accuracy: 0.8626

Epoch 16/20



```

30/30 [=====] - 1s 18ms/step - loss: 0.0174 - accuracy:
0.9977 - val_loss: 0.5550 - val_accuracy: 0.8685
Epoch 17/20
30/30 [=====] - 1s 21ms/step - loss: 0.0147 - accuracy:
0.9978 - val_loss: 0.5917 - val_accuracy: 0.8649
Epoch 18/20
30/30 [=====] - 1s 19ms/step - loss: 0.0108 - accuracy:
0.9992 - val_loss: 0.6251 - val_accuracy: 0.8672
Epoch 19/20
30/30 [=====] - 1s 20ms/step - loss: 0.0071 - accuracy:
0.9997 - val_loss: 0.6602 - val_accuracy: 0.8663
Epoch 20/20
30/30 [=====] - 1s 21ms/step - loss: 0.0057 - accuracy:
0.9997 - val_loss: 0.7050 - val_accuracy: 0.8661

```

```

[13]: history_dict = history.history
      history_dict.keys()

```

```

[13]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```

### Plotting the training and validation loss

```

[15]: import matplotlib.pyplot as plt


      history_dict = history.history
      loss_values = history_dict['loss']
      val_loss_values = history_dict['val_loss']

      epochs = range(1, len(loss_values) + 1)

      plt.plot(epochs, loss_values, 'bo', label='Training loss')
      plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()

      plt.show()

```




#### Plotting the training and validation accuracy

```
[18]: plt.clf()
      acc = history_dict['accuracy']
      val_acc = history_dict['val_accuracy']

      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()

      plt.show()
```



### Retraining a model from scratch

```
[21]: model = keras.Sequential([
    layers.Dense(16, activation='relu'),
    layers.Dense(16, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

Epoch 1/4

49/49 [=====] - 2s 13ms/step - loss: 0.5546 - accuracy: 0.7265

Epoch 2/4

49/49 [=====] - 1s 13ms/step - loss: 0.2702 - accuracy: 0.9123

Epoch 3/4

49/49 [=====] - 1s 14ms/step - loss: 0.1997 - accuracy: 0.9313

Epoch 4/4

```
49/49 [=====] - 1s 13ms/step - loss: 0.1605 - accuracy: 0.9444
782/782 [=====] - 2s 2ms/step - loss: 0.3289 - accuracy: 0.8707: 0s - loss:
```

```
[22]: results
```

```
[22]: [0.3288571238517761, 0.8706799745559692]
```

### Using a trained network to generate predictions on new data

```
[23]: model.predict(x_test)
```

```
[23]: array([[0.22420892],
             [0.9990783 ],
             [0.9645244 ],
             ...,
             [0.185489  ],
             [0.12515825],
             [0.80406165]], dtype=float32)
```

## 0.2 Assignment 5.2 -

Classifying newswires: a multiclass classification example Implement the news classifier found in section 3.5 of Deep Learning with Python.

### Load data

```
[1]: from keras.datasets import reuters

(train_data, train_labels), (test_data, test_labels) = reuters.load_data(
    num_words=10000)
```

```
C:\Users\Tushar\AppData\Roaming\Python\Python38\site-
packages\tensorflow\python\keras\datasets\reuters.py:148:
VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences
(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths
or shapes) is deprecated. If you meant to do this, you must specify
'dtype=object' when creating the ndarray
    x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])
C:\Users\Tushar\AppData\Roaming\Python\Python38\site-
packages\tensorflow\python\keras\datasets\reuters.py:149:
VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences
(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths
or shapes) is deprecated. If you meant to do this, you must specify
'dtype=object' when creating the ndarray
    x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

```
[2]: len(train_data)
```

```
[2]: 8982
```

```
[3]: len(test_data)
```

```
[3]: 2246
```

```
[4]: train_data[10]
```

```
[4]: [1,  
      245,  
      273,  
      207,  
      156,  
      53,  
      74,  
      160,  
      26,  
      14,  
      46,  
      296,  
      26,  
      39,  
      74,  
      2979,  
      3554,  
      14,  
      46,  
      4689,  
      4329,  
      86,  
      61,  
      3499,  
      4795,  
      14,  
      61,  
      451,  
      4329,  
      17,  
      12]
```

#### Decoding newswires back to text

```
[5]: word_index = reuters.get_word_index()  
      reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])  
      decoded_newswire = ' '.join([reverse_word_index.get(i - 3, '?') for i in  
                                    train_data[0]])
```

```
[6]: train_labels[10]
```

[6]: 3

### Preparing the data

```
[7]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
[8]: def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

one_hot_train_labels = to_one_hot(train_labels)
one_hot_test_labels = to_one_hot(test_labels)
```

```
[9]: from keras.utils.np_utils import to_categorical

one_hot_train_labels = to_categorical(train_labels)
one_hot_test_labels = to_categorical(test_labels)
```

### Building your network

```
[10]: from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
```

### Compiling the model

```
[11]: model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

### Validating your approach

```
[12]: x_val = x_train[:1000]
partial_x_train = x_train[1000:]
```

```
y_val = one_hot_train_labels[:1000]
partial_y_train = one_hot_train_labels[1000:]
```

### Training the model

```
[13]: history = model.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))
```

Epoch 1/20

16/16 [=====] - 8s 206ms/step - loss: 3.2135 - accuracy: 0.3811 - val\_loss: 1.8512 - val\_accuracy: 0.6290

Epoch 2/20

16/16 [=====] - 1s 39ms/step - loss: 1.6128 - accuracy: 0.6666 - val\_loss: 1.3601 - val\_accuracy: 0.7090

Epoch 3/20

16/16 [=====] - 1s 41ms/step - loss: 1.1183 - accuracy: 0.7668 - val\_loss: 1.1756 - val\_accuracy: 0.7570

Epoch 4/20

16/16 [=====] - 1s 46ms/step - loss: 0.9017 - accuracy: 0.8116 - val\_loss: 1.0637 - val\_accuracy: 0.7770

Epoch 5/20

16/16 [=====] - 1s 45ms/step - loss: 0.6961 - accuracy: 0.8586 - val\_loss: 0.9799 - val\_accuracy: 0.7990

Epoch 6/20

16/16 [=====] - 1s 42ms/step - loss: 0.5594 - accuracy: 0.8881 - val\_loss: 0.9449 - val\_accuracy: 0.8020

Epoch 7/20

16/16 [=====] - 1s 38ms/step - loss: 0.4525 - accuracy: 0.9135 - val\_loss: 0.9265 - val\_accuracy: 0.8020

Epoch 8/20

16/16 [=====] - 1s 39ms/step - loss: 0.3553 - accuracy: 0.9325 - val\_loss: 0.8982 - val\_accuracy: 0.8170

Epoch 9/20

16/16 [=====] - 1s 37ms/step - loss: 0.3046 - accuracy: 0.9385 - val\_loss: 0.9095 - val\_accuracy: 0.8150

Epoch 10/20

16/16 [=====] - 1s 37ms/step - loss: 0.2458 - accuracy: 0.9472 - val\_loss: 0.9551 - val\_accuracy: 0.8010

Epoch 11/20

16/16 [=====] - 1s 37ms/step - loss: 0.2092 - accuracy: 0.9552 - val\_loss: 0.9322 - val\_accuracy: 0.8200

Epoch 12/20

16/16 [=====] - 1s 33ms/step - loss: 0.1874 - accuracy: 0.9561 - val\_loss: 0.9483 - val\_accuracy: 0.8100

Epoch 13/20

```

16/16 [=====] - 1s 34ms/step - loss: 0.1601 - accuracy:
0.9601 - val_loss: 0.9672 - val_accuracy: 0.8150
Epoch 14/20
16/16 [=====] - 1s 39ms/step - loss: 0.1451 - accuracy:
0.9595 - val_loss: 1.0671 - val_accuracy: 0.8050
Epoch 15/20
16/16 [=====] - 1s 36ms/step - loss: 0.1327 - accuracy:
0.9599 - val_loss: 1.0189 - val_accuracy: 0.8110
Epoch 16/20
16/16 [=====] - 1s 35ms/step - loss: 0.1259 - accuracy:
0.9619 - val_loss: 1.0720 - val_accuracy: 0.8030
Epoch 17/20
16/16 [=====] - 1s 35ms/step - loss: 0.1163 - accuracy:
0.9626 - val_loss: 1.0739 - val_accuracy: 0.8020
Epoch 18/20
16/16 [=====] - 1s 38ms/step - loss: 0.1099 - accuracy:
0.9602 - val_loss: 1.0494 - val_accuracy: 0.8130
Epoch 19/20
16/16 [=====] - 1s 35ms/step - loss: 0.1023 - accuracy:
0.9639 - val_loss: 1.1001 - val_accuracy: 0.8000
Epoch 20/20
16/16 [=====] - 1s 36ms/step - loss: 0.0996 - accuracy:
0.9656 - val_loss: 1.1330 - val_accuracy: 0.8070

```

### Plotting the training and validation loss

```

[14]: import matplotlib.pyplot as plt

loss = history.history['loss']
val_loss = history.history['val_loss']


epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```






#### Plotting the training and validation accuracy

```
[16]: plt.clf()

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



### Retraining a model from scratch

```
[17]: model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(partial_x_train,
        partial_y_train,
        epochs=9,
        batch_size=512,
        validation_data=(x_val, y_val))
results = model.evaluate(x_test, one_hot_test_labels)
```

Epoch 1/9

16/16 [=====] - 2s 66ms/step - loss: 3.0504 - accuracy: 0.4232 - val\_loss: 1.7205 - val\_accuracy: 0.6410

Epoch 2/9

16/16 [=====] - 1s 41ms/step - loss: 1.5002 - accuracy: 0.6956 - val\_loss: 1.3097 - val\_accuracy: 0.7100

Epoch 3/9

16/16 [=====] - 1s 45ms/step - loss: 1.0782 - accuracy:

```

0.7692 - val_loss: 1.1295 - val_accuracy: 0.7500
Epoch 4/9
16/16 [=====] - 1s 46ms/step - loss: 0.8361 - accuracy:
0.8212 - val_loss: 1.0276 - val_accuracy: 0.7800
Epoch 5/9
16/16 [=====] - 1s 47ms/step - loss: 0.6682 - accuracy:
0.8611 - val_loss: 0.9737 - val_accuracy: 0.7910
Epoch 6/9
16/16 [=====] - 1s 49ms/step - loss: 0.5204 - accuracy:
0.8938 - val_loss: 0.9255 - val_accuracy: 0.8090
Epoch 7/9
16/16 [=====] - ETA: 0s - loss: 0.4221 - accuracy: 0.91
- 1s 46ms/step - loss: 0.4214 - accuracy: 0.9141 - val_loss: 0.9039 -
val_accuracy: 0.8070
Epoch 8/9
16/16 [=====] - 1s 77ms/step - loss: 0.3331 - accuracy:
0.9286 - val_loss: 0.9156 - val_accuracy: 0.8090
Epoch 9/9
16/16 [=====] - 1s 48ms/step - loss: 0.2814 - accuracy:
0.9378 - val_loss: 0.8833 - val_accuracy: 0.8150
71/71 [=====] - 0s 5ms/step - loss: 0.9773 - accuracy:
0.7832

```

```
[18]: results
```

```
[18]: [0.9773288369178772, 0.7831701040267944]
```

```
[20]: import copy
test_labels_copy = copy.copy(test_labels)
np.random.shuffle(test_labels_copy)
hits_array = np.array(test_labels) == np.array(test_labels_copy)
float(np.sum(hits_array)) / len(test_labels)
```

```
[20]: 0.19412288512911843
```

### Generating predictions on new data

```
[21]: predictions = model.predict(x_test)
```

```
[22]: predictions[0].shape
```

```
[22]: (46,)
```

```
[23]: np.sum(predictions[0])
```

```
[23]: 1.0
```

```
[24]: np.argmax(predictions[0])
```

[24]: 3

#### A different way to handle the labels and the loss

```
[25]: y_train = np.array(train_labels)
      y_test = np.array(test_labels)
```

```
[26]: model.compile(optimizer='rmsprop',
                    loss='sparse_categorical_crossentropy',
                    metrics=['acc'])
```

#### The importance of having sufficiently large intermediate layers

```
[27]: model = models.Sequential()
      model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
      model.add(layers.Dense(4, activation='relu'))
      model.add(layers.Dense(46, activation='softmax'))

      model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
      model.fit(partial_x_train,
                partial_y_train,
                epochs=20,
                batch_size=128,
                validation_data=(x_val, y_val))
```

Epoch 1/20

63/63 [=====] - 3s 20ms/step - loss: 3.6646 - accuracy: 0.0873 - val\_loss: 3.1689 - val\_accuracy: 0.2580

Epoch 2/20

63/63 [=====] - 1s 14ms/step - loss: 2.9481 - accuracy: 0.2650 - val\_loss: 2.5934 - val\_accuracy: 0.2940

Epoch 3/20

63/63 [=====] - 1s 16ms/step - loss: 2.3941 - accuracy: 0.3167 - val\_loss: 2.1410 - val\_accuracy: 0.3460

Epoch 4/20

63/63 [=====] - 1s 13ms/step - loss: 1.8790 - accuracy: 0.3711 - val\_loss: 1.7268 - val\_accuracy: 0.4880

Epoch 5/20

63/63 [=====] - 1s 15ms/step - loss: 1.5390 - accuracy: 0.5170 - val\_loss: 1.6470 - val\_accuracy: 0.4750

Epoch 6/20

63/63 [=====] - 5s 73ms/step - loss: 1.4122 - accuracy: 0.5441 - val\_loss: 1.5302 - val\_accuracy: 0.6410

Epoch 7/20

63/63 [=====] - 1s 20ms/step - loss: 1.2311 - accuracy: 0.6838 - val\_loss: 1.4330 - val\_accuracy: 0.6650

Epoch 8/20

```

63/63 [=====] - 1s 19ms/step - loss: 1.0729 - accuracy:
0.7224 - val_loss: 1.4012 - val_accuracy: 0.6650
Epoch 9/20
63/63 [=====] - 1s 23ms/step - loss: 0.9657 - accuracy:
0.7502 - val_loss: 1.4317 - val_accuracy: 0.6750
Epoch 10/20
63/63 [=====] - 1s 22ms/step - loss: 0.9137 - accuracy:
0.7533 - val_loss: 1.4498 - val_accuracy: 0.6730
Epoch 11/20
63/63 [=====] - 1s 15ms/step - loss: 0.8862 - accuracy:
0.7634 - val_loss: 1.4923 - val_accuracy: 0.6750
Epoch 12/20
63/63 [=====] - 1s 15ms/step - loss: 0.8291 - accuracy:
0.7802 - val_loss: 1.5320 - val_accuracy: 0.6690
Epoch 13/20
63/63 [=====] - 1s 13ms/step - loss: 0.8111 - accuracy:
0.7818 - val_loss: 1.5888 - val_accuracy: 0.6700
Epoch 14/20
63/63 [=====] - 1s 15ms/step - loss: 0.7676 - accuracy:
0.7955 - val_loss: 1.5891 - val_accuracy: 0.6790
Epoch 15/20
63/63 [=====] - 1s 14ms/step - loss: 0.7723 - accuracy:
0.7937 - val_loss: 1.6605 - val_accuracy: 0.6720
Epoch 16/20
63/63 [=====] - 1s 15ms/step - loss: 0.7140 - accuracy:
0.8067 - val_loss: 1.7169 - val_accuracy: 0.6750
Epoch 17/20
63/63 [=====] - 1s 16ms/step - loss: 0.7127 - accuracy:
0.8121 - val_loss: 1.7745 - val_accuracy: 0.6740
Epoch 18/20
63/63 [=====] - 1s 14ms/step - loss: 0.6930 - accuracy:
0.8142 - val_loss: 1.8306 - val_accuracy: 0.6800
Epoch 19/20
63/63 [=====] - 1s 15ms/step - loss: 0.6749 - accuracy:
0.8178 - val_loss: 1.8679 - val_accuracy: 0.6780
Epoch 20/20
63/63 [=====] - 1s 14ms/step - loss: 0.6562 - accuracy:
0.8202 - val_loss: 1.9090 - val_accuracy: 0.6770

```

[27]: <tensorflow.python.keras.callbacks.History at 0x25786e132b0>

### 0.3 Assignment 5.3

Implement the housing price regression model found in section 3.6 of Deep Learning with Python.

#### The Boston Housing Price dataset

```
[28]: from keras.datasets import boston_housing
```

```
(train_data, train_targets), (test_data, test_targets) = boston_housing.  
    ↪load_data()
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston\\_housing.npz](https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston_housing.npz)

57344/57026 [=====] - 0s 1us/step

```
[29]: train_data.shape
```

```
[29]: (404, 13)
```

```
[30]: test_data.shape
```

```
[30]: (102, 13)
```

```
[31]: train_targets
```

```
[31]: array([15.2, 42.3, 50. , 21.1, 17.7, 18.5, 11.3, 15.6, 15.6, 14.4, 12.1,  
        17.9, 23.1, 19.9, 15.7,  8.8, 50. , 22.5, 24.1, 27.5, 10.9, 30.8,  
        32.9, 24. , 18.5, 13.3, 22.9, 34.7, 16.6, 17.5, 22.3, 16.1, 14.9,  
        23.1, 34.9, 25. , 13.9, 13.1, 20.4, 20. , 15.2, 24.7, 22.2, 16.7,  
        12.7, 15.6, 18.4, 21. , 30.1, 15.1, 18.7,  9.6, 31.5, 24.8, 19.1,  
        22. , 14.5, 11. , 32. , 29.4, 20.3, 24.4, 14.6, 19.5, 14.1, 14.3,  
        15.6, 10.5,  6.3, 19.3, 19.3, 13.4, 36.4, 17.8, 13.5, 16.5,  8.3,  
        14.3, 16. , 13.4, 28.6, 43.5, 20.2, 22. , 23. , 20.7, 12.5, 48.5,  
        14.6, 13.4, 23.7, 50. , 21.7, 39.8, 38.7, 22.2, 34.9, 22.5, 31.1,  
        28.7, 46. , 41.7, 21. , 26.6, 15. , 24.4, 13.3, 21.2, 11.7, 21.7,  
        19.4, 50. , 22.8, 19.7, 24.7, 36.2, 14.2, 18.9, 18.3, 20.6, 24.6,  
        18.2,  8.7, 44. , 10.4, 13.2, 21.2, 37. , 30.7, 22.9, 20. , 19.3,  
        31.7, 32. , 23.1, 18.8, 10.9, 50. , 19.6,  5. , 14.4, 19.8, 13.8,  
        19.6, 23.9, 24.5, 25. , 19.9, 17.2, 24.6, 13.5, 26.6, 21.4, 11.9,  
        22.6, 19.6,  8.5, 23.7, 23.1, 22.4, 20.5, 23.6, 18.4, 35.2, 23.1,  
        27.9, 20.6, 23.7, 28. , 13.6, 27.1, 23.6, 20.6, 18.2, 21.7, 17.1,  
        8.4, 25.3, 13.8, 22.2, 18.4, 20.7, 31.6, 30.5, 20.3,  8.8, 19.2,  
        19.4, 23.1, 23. , 14.8, 48.8, 22.6, 33.4, 21.1, 13.6, 32.2, 13.1,  
        23.4, 18.9, 23.9, 11.8, 23.3, 22.8, 19.6, 16.7, 13.4, 22.2, 20.4,  
        21.8, 26.4, 14.9, 24.1, 23.8, 12.3, 29.1, 21. , 19.5, 23.3, 23.8,  
        17.8, 11.5, 21.7, 19.9, 25. , 33.4, 28.5, 21.4, 24.3, 27.5, 33.1,  
        16.2, 23.3, 48.3, 22.9, 22.8, 13.1, 12.7, 22.6, 15. , 15.3, 10.5,  
        24. , 18.5, 21.7, 19.5, 33.2, 23.2,  5. , 19.1, 12.7, 22.3, 10.2,  
        13.9, 16.3, 17. , 20.1, 29.9, 17.2, 37.3, 45.4, 17.8, 23.2, 29. ,  
        22. , 18. , 17.4, 34.6, 20.1, 25. , 15.6, 24.8, 28.2, 21.2, 21.4,  
        23.8, 31. , 26.2, 17.4, 37.9, 17.5, 20. ,  8.3, 23.9,  8.4, 13.8,  
        7.2, 11.7, 17.1, 21.6, 50. , 16.1, 20.4, 20.6, 21.4, 20.6, 36.5,  
        8.5, 24.8, 10.8, 21.9, 17.3, 18.9, 36.2, 14.9, 18.2, 33.3, 21.8,  
        19.7, 31.6, 24.8, 19.4, 22.8,  7.5, 44.8, 16.8, 18.7, 50. , 50. ,
```

```
19.5, 20.1, 50. , 17.2, 20.8, 19.3, 41.3, 20.4, 20.5, 13.8, 16.5,
23.9, 20.6, 31.5, 23.3, 16.8, 14. , 33.8, 36.1, 12.8, 18.3, 18.7,
19.1, 29. , 30.1, 50. , 50. , 22. , 11.9, 37.6, 50. , 22.7, 20.8,
23.5, 27.9, 50. , 19.3, 23.9, 22.6, 15.2, 21.7, 19.2, 43.8, 20.3,
33.2, 19.9, 22.5, 32.7, 22. , 17.1, 19. , 15. , 16.1, 25.1, 23.7,
28.7, 37.2, 22.6, 16.4, 25. , 29.8, 22.1, 17.4, 18.1, 30.3, 17.5,
24.7, 12.6, 26.5, 28.7, 13.3, 10.4, 24.4, 23. , 20. , 17.8, 7. ,
11.8, 24.4, 13.8, 19.4, 25.2, 19.4, 19.4, 29.1])
```

### Preparing the data

```
[32]: mean = train_data.mean(axis=0)
train_data -= mean
std = train_data.std(axis=0)
train_data /= std

test_data -= mean
test_data /= std
```

### Building your network

```
[33]: import keras
from keras import models
from keras import layers

def build_model():
    model = models.Sequential()
    model.add(layers.Dense(64, activation='relu',
                           input_shape=(train_data.shape[1],)))
    model.add(layers.Dense(64, activation='relu'))
    model.add(layers.Dense(1))
    model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
    return model
```

### Validating your approach using K-fold validation

```
[35]: import numpy as np

k = 4
num_val_samples = len(train_data) // k
num_epochs = 100
all_scores = []
for i in range(k):
    print('processing fold #', i)
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]

    partial_train_data = np.concatenate(
        [train_data[:i * num_val_samples],
```

```

        train_data[(i + 1) * num_val_samples:],
        axis=0)
    partial_train_targets = np.concatenate(
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]],
        axis=0)

    model = build_model()
    model.fit(partial_train_data, partial_train_targets,
              epochs=num_epochs, batch_size=1, verbose=0)
    val_mse, val_mae = model.evaluate(val_data, val_targets, verbose=0)
    all_scores.append(val_mae)

```

```

processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3

```

```
[36]: all_scores
```

```
[36]: [1.982920527458191, 2.484312057495117, 2.6799979209899902, 2.448901891708374]
```

```
[37]: np.mean(all_scores)
```

```
[37]: 2.399033099412918
```

### Saving the validation logs at each fold

```

[39]: num_epochs = 500
      all_mae_histories = []
      for i in range(k):
          print('processing fold #', i)
          val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
          val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
          partial_train_data = np.concatenate(
              [train_data[:i * num_val_samples],
               train_data[(i + 1) * num_val_samples:]],
              axis=0)
          partial_train_targets = np.concatenate(
              [train_targets[:i * num_val_samples],
               train_targets[(i + 1) * num_val_samples:]],
              axis=0)

          model = build_model()
          history = model.fit(partial_train_data, partial_train_targets,
                              validation_data=(val_data, val_targets),
                              epochs=num_epochs, batch_size=1, verbose=0)
          mae_history = history.history['val_mae']
          all_mae_histories.append(mae_history)

```



```
processing fold # 0
processing fold # 1
processing fold # 2
processing fold # 3
```


#### Building the history of successive mean K-fold validation scores

```
[40]: average_mae_history = [
        np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)]
```

#### Plotting validation scores

```
[41]: import matplotlib.pyplot as plt

plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()
```



#### Plotting validation scores, excluding the first 10 data points

```
[42]: def smooth_curve(points, factor=0.9):
        smoothed_points = []
        for point in points:
            if smoothed_points:
                previous = smoothed_points[-1]
                smoothed_points.append(previous * factor + point * (1 - factor))
            else:
                smoothed_points.append(point)
```


```

    else:
        smoothed_points.append(point)
    return smoothed_points

smooth_mae_history = smooth_curve(average_mae_history[10:])

plt.plot(range(1, len(smooth_mae_history) + 1), smooth_mae_history)
plt.xlabel('Epochs')
plt.ylabel('Validation MAE')
plt.show()

```



### Training the final model

```

[44]: model = build_model()
      model.fit(train_data, train_targets,
                epochs=80, batch_size=16, verbose=0)
      test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)

```

4/4 [=====] - 0s 3ms/step - loss: 16.6004 - mae: 2.6129

```

[45]: test_mae_score

```

```

[45]: 2.612933874130249

```

### Generating predictions on new data

```
[46]: predictions = model.predict(test_data)
      predictions[0]
```

```
[46]: array([8.1121025], dtype=float32)
```

```
[ ]:
```