

assignment06-1_muley_tushar

January 8, 2022

Name: Tushar Muley

Assignment: Assignment 6-1

Date: January 9, 2022

Assignment 6.1 Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a ConvNet model that classifies images in the MNIST digit dataset.

```
[1]: from keras import layers
     from keras import models
```

```
[2]: model = models.Sequential()
     model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
     model.add(layers.MaxPooling2D((2, 2)))
     model.add(layers.Conv2D(64, (3, 3), activation='relu'))
     model.add(layers.MaxPooling2D((2, 2)))
     model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
[3]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
Total params: 55,744		
Trainable params: 55,744		
Non-trainable params: 0		

Listing 5.2. Adding a classifier on top of the convnet

```
[4]: model.add(layers.Flatten())
      model.add(layers.Dense(64, activation='relu'))
      model.add(layers.Dense(10, activation='softmax'))
```

```
[5]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

Listing 5.3. Training the convnet on MNIST images

```
[6]: from keras.datasets import mnist
      from keras.utils import to_categorical
```

```
[7]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
[8]: train_images = train_images.reshape((60000, 28, 28, 1))
      train_images = train_images.astype('float32') / 255
```

```
[9]: test_images = test_images.reshape((10000, 28, 28, 1))
      test_images = test_images.astype('float32') / 255
```

```
[10]: train_labels = to_categorical(train_labels)
       test_labels = to_categorical(test_labels)
```

```
[12]: model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

```
[16]: # validation set
val_images = train_images[:10000,:]
val_labels = train_labels[:10000,:]

partial_train_images = train_images[10000:,:]
partial_train_labels = train_labels[10000:,:]

partial_train_images.shape, val_images.shape
```

```
[16]: ((50000, 28, 28, 1), (10000, 28, 28, 1))
```

```
[18]: # train the model
history=model.fit(train_images,
                  train_labels,
                  epochs=5,
                  batch_size=64,
                  validation_data=(val_images, val_labels))
```

```
Epoch 1/5
938/938 [=====] - 35s 36ms/step - loss: 0.0154 -
accuracy: 0.9953 - val_loss: 0.0124 - val_accuracy: 0.9962
Epoch 2/5
938/938 [=====] - 44s 47ms/step - loss: 0.0128 -
accuracy: 0.9964 - val_loss: 0.0080 - val_accuracy: 0.9974
Epoch 3/5
938/938 [=====] - 45s 48ms/step - loss: 0.0100 -
accuracy: 0.9970 - val_loss: 0.0045 - val_accuracy: 0.9987
Epoch 4/5
938/938 [=====] - 43s 46ms/step - loss: 0.0077 -
accuracy: 0.9977 - val_loss: 0.0045 - val_accuracy: 0.9985
Epoch 5/5
938/938 [=====] - 40s 43ms/step - loss: 0.0076 -
accuracy: 0.9977 - val_loss: 0.0068 - val_accuracy: 0.9976
```


```
[19]: history_dict = history.history
      history_dict.keys()
```

```
[19]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

Plot

Training and Validation Loss


```
[21]: import matplotlib.pyplot as plt
history_dict = history.history
loss_val = history_dict["loss"]
val_loss_val = history_dict["val_loss"]
epochs= range(1, len(loss_val) +1)
plt.plot(epochs, loss_val, "bo", label="Training Loss")
plt.plot(epochs, val_loss_val, "b", label="Validation Loss")
plt.plot(title = "Training and Validation Loss")
plt.plot(xlabel="Epochs")
plt.plot(ylabel="Loss")
plt.legend()
plt.show()
```



Training and Validation Accuracy

```
[22]: history_dict = history.history
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
epochs= range(1, len(loss_val) +1)
plt.plot(epochs, acc, "bo", label="Training Accuracy")
plt.plot(epochs, val_acc, "b", label="Validation Accuracy")
plt.plot(title = "Training and Validation Accuracy")
plt.plot(xlabel="Epochs")
plt.plot(ylabel="Accuracy")
plt.legend()
```

```
plt.show()
```



Retrain model

```
[26]: model.compile(optimizer='rmsprop',  
                    loss = 'categorical_crossentropy',  
                    metrics = ['accuracy'])  
  
model.fit(train_images, train_labels, epochs=3, batch_size=512)  
  
results=model.evaluate(test_images, test_labels)
```

Epoch 1/3

118/118 [=====] - 30s 248ms/step - loss: 0.0018 -
accuracy: 0.9993

Epoch 2/3

118/118 [=====] - 34s 283ms/step - loss: 6.9039e-04 -
accuracy: 0.9997

Epoch 3/3

118/118 [=====] - 33s 276ms/step - loss: 5.5668e-04 -
accuracy: 0.9998

313/313 [=====] - 3s 8ms/step - loss: 0.0373 -
accuracy: 0.9927

```
[27]: results
```

```
[27]: [0.03726635128259659, 0.9926999807357788]
```

```
[28]: model.predict(test_images)
```

```
[28]: array([[1.17138473e-18, 1.52956668e-14, 1.89267655e-13, ...,  
            1.00000000e+00, 2.82987699e-16, 9.99412358e-13],  
            [3.89941445e-19, 2.53855212e-17, 1.00000000e+00, ...,  
            6.27537656e-27, 4.38832692e-20, 1.33996726e-29],  
            [5.93215256e-13, 1.00000000e+00, 1.57812270e-11, ...,  
            8.91156648e-10, 1.22106908e-10, 1.01484397e-10],  
            ...,  
            [2.81409270e-33, 1.80594282e-24, 6.73015316e-29, ...,  
            2.94729163e-23, 3.30883150e-21, 8.33329335e-22],  
            [5.74947273e-21, 1.00544685e-23, 1.31905212e-29, ...,  
            4.42883865e-28, 1.92395280e-13, 1.18429775e-18],  
            [3.33153792e-15, 6.87902423e-18, 2.37463993e-14, ...,  
            9.12308501e-30, 2.69159071e-13, 1.71166226e-21]], dtype=float32)
```

```
[ ]:
```