

Assignment10_Muley_Tushar

January 25, 2022

Name: Tushar Muley

Assignment: Assignment 10

Date: January 30, 2022

0.0.1 Assignment 10.1

Assignment 10.1a

```
[1]: # import libraries
import string
import numpy as np
```

```
[2]: # the tokenizer fuction
def tokenize(sentence):
    tokens = []
    # remove punctuations
    sentence = sentence.translate(str.maketrans('', '', string.punctuation))
    tokens = sentence.split()
    return tokens
```

```
[4]: # check the tokenizer function
sentence = "Hello! Welcome to the class! This will be your last class before_
→you graduate."
tokens = tokenize(sentence)
print(tokens)
```

```
['Hello', 'Welcome', 'to', 'the', 'class', 'This', 'will', 'be', 'your', 'last',
'class', 'before', 'you', 'graduate']
```

Assignment 10.1b

```
[5]: # ngrams fuction
def ngrams(tokens, n):
    ngrams = []
    for i in range(len(tokens) - n + 1):
        ngram = []
        for a in range(n):
            ngram.append(tokens[i+a])
```

```

    ngrams.append(ngram)
return ngrams

```

```

[6]: # check the ngrams function
ngrams = ngrams(tokens, 3)
print(ngrams)

```

```

[['Hello', 'Welcome', 'to'], ['Welcome', 'to', 'the'], ['to', 'the', 'class'],
['the', 'class', 'This'], ['class', 'This', 'will'], ['This', 'will', 'be'],
['will', 'be', 'your'], ['be', 'your', 'last'], ['your', 'last', 'class'],
['last', 'class', 'before'], ['class', 'before', 'you'], ['before', 'you',
'graduate']]

```

Assignment 10.1c

```

[7]: # one_hot_encode function

def one_hot_encode(tokens, num_words):
    token_index = {}
    for token in tokens:
        if token not in token_index:
            token_index[token] = len(token_index) + 1

    results = np.zeros(shape=(num_words, max(token_index.values()) + 1))

    for i, token in list(enumerate(tokens))[:num_words]:
        index = token_index.get(token)
        results[i, index] = 1.

    return results

```

```

[8]: # check the one_hot_encode function

results = one_hot_encode(tokens, 20)
print(results)

```

```

[[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]

```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

Assignment 10.2 Using listings 6.16, 6.17, and 6.18 in Deep Learning with Python as a guide, train a sequential model with embeddings on the IMDB data found in `data/external/imdb/`. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

```
[36]: # import libraries
import os, pathlib, shutil, random
from pathlib import Path
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
import numpy as np
import matplotlib.pyplot as plt
from keras.layers import LSTM
from keras import layers
from keras.optimizers import RMSprop
```

```
[10]: !curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -xf aclImdb_v1.tar.gz
```

	% Total		% Received		% Xferd		Average Speed		Time		Time		Time		Current
							Dload Upload		Total		Spent		Left		Speed
100	80.2M	100	80.2M	0	0	2884k	0	0:00:28	0:00:28	--:--:--	4590k				

```
[11]: !rm -r aclImdb/train/unsup
```

```
[12]: !cat aclImdb/train/pos/4077_10.txt
```

I first saw this back in the early 90s on UK TV, i did like it then but i missed the chance to tape it, many years passed but the film always stuck with me and i lost hope of seeing it TV again, the main thing that stuck with me was the end, the hole castle part really touched me, its easy to watch, has a great story, great music, the list goes on and on, its OK me saying how good it is but everyone will take there own best bits away with them once they have seen it, yes the animation is top notch and beautiful to watch, it does show its age in a very few parts but that has now become part of it beauty, i am so glad it has came out on DVD as it is one of my top 10 films of all time. Buy it or rent it just see it, best viewing is at night alone with drink and food in reach so you

don't have to stop the film.

Enjoy

```
[13]: base_dir = pathlib.Path("aclImdb")
      val_dir = base_dir / "val"
      train_dir = base_dir / "train"
      for category in ("neg", "pos"):
          os.makedirs(val_dir / category)
          files = os.listdir(train_dir / category)
          random.Random(1337).shuffle(files)
          num_val_samples = int(0.2 * len(files))
          val_files = files[-num_val_samples:]
          for fname in val_files:
              shutil.move(train_dir / category / fname, val_dir / category / fname)
```

```
[14]: current_dir = Path(os.getcwd()).absolute()
      imdb_dir = current_dir.joinpath('aclImdb')
      train_dir = os.path.join(imdb_dir, 'train')
```

```
[15]: labels = []
      texts = []

      for label_type in ['neg', 'pos']:
          dir_name = os.path.join(train_dir, label_type)
          for fname in os.listdir(dir_name):
              if fname[-4:] == '.txt':
                  f = open(os.path.join(dir_name, fname))
                  texts.append(f.read())
                  f.close()
                  if label_type == 'neg':
                      labels.append(0)
                  else:
                      labels.append(1)
```

Tokenizing the text of the IMDB data

```
[16]: maxlen = 100
      training_samples = 200
      validation_samples = 10000
      max_words = 10000
      embedding_dim = 100
```

```
[17]: tokenizer = Tokenizer(num_words=max_words)
      tokenizer.fit_on_texts(texts)
      sequences = tokenizer.texts_to_sequences(texts)
```

```
[18]: word_index = tokenizer.word_index
      print('Found %s unique tokens' % len(word_index))
```

Found 80258 unique tokens

```
[19]: data = pad_sequences(sequences, maxlen=maxlen)
```

```
[20]: labels = np.asarray(labels)
      print('Shape of data tensor:', data.shape)
      print('Shape of label tensor:', labels.shape)
```

Shape of data tensor: (20000, 100)

Shape of label tensor: (20000,)

```
[21]: indices = np.arange(data.shape[0])
      np.random.shuffle(indices)
      data = data[indices]
      labels = labels[indices]
```

```
[22]: # split the data
      x_train = data[:training_samples]
      y_train = labels[:training_samples]
      x_val = data[training_samples: training_samples + validation_samples]
      y_val = labels[training_samples: training_samples + validation_samples]
```

Train the model

```
[23]: model = Sequential()
      model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
      model.add(Flatten())
      model.add(Dense(32, activation='relu'))
      model.add(Dense(1, activation='sigmoid'))
      model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	1000000
flatten (Flatten)	(None, 10000)	0
dense (Dense)	(None, 32)	320032
dense_1 (Dense)	(None, 1)	33

=====
Total params: 1,320,065
Trainable params: 1,320,065
Non-trainable params: 0
=====

```
[24]: model.compile(optimizer='rmsprop',loss='binary_crossentropy',metrics=['acc'])
```

```
[25]: history = model.fit(x_train, y_train,
    ↪ epochs=10,batch_size=32,validation_data=(x_val, y_val))
```


```
Epoch 1/10
7/7 [=====] - 1s 130ms/step - loss: 0.6952 - acc:
0.5100 - val_loss: 0.6934 - val_acc: 0.5130
Epoch 2/10
7/7 [=====] - 1s 107ms/step - loss: 0.5170 - acc:
0.9850 - val_loss: 0.7009 - val_acc: 0.5103
Epoch 3/10
7/7 [=====] - 1s 115ms/step - loss: 0.3082 - acc:
0.9950 - val_loss: 0.6978 - val_acc: 0.5142
Epoch 4/10
7/7 [=====] - 1s 114ms/step - loss: 0.1430 - acc:
1.0000 - val_loss: 0.7186 - val_acc: 0.5084
Epoch 5/10
7/7 [=====] - 1s 111ms/step - loss: 0.0681 - acc:
1.0000 - val_loss: 0.6968 - val_acc: 0.5342
Epoch 6/10
7/7 [=====] - 1s 109ms/step - loss: 0.0345 - acc:
1.0000 - val_loss: 0.7414 - val_acc: 0.5150
Epoch 7/10
7/7 [=====] - 1s 107ms/step - loss: 0.0194 - acc:
1.0000 - val_loss: 0.7112 - val_acc: 0.5260
Epoch 8/10
7/7 [=====] - 1s 108ms/step - loss: 0.0105 - acc:
1.0000 - val_loss: 0.7105 - val_acc: 0.5332
Epoch 9/10
7/7 [=====] - 1s 107ms/step - loss: 0.0063 - acc:
1.0000 - val_loss: 0.7206 - val_acc: 0.5265
Epoch 10/10
7/7 [=====] - 1s 110ms/step - loss: 0.0039 - acc:
1.0000 - val_loss: 0.7252 - val_acc: 0.5285
```


```
[26]: # plot the data
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and Validation of Accuracy')
plt.legend()
```

```
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training Loss')
plt.plot(epochs, val_loss, 'b', label='Validation Loss')
plt.title('Training and Validation of Loss')
plt.legend()

plt.show()
```





The Test Data Set

```
[27]: test_dir = os.path.join(imdb_dir, 'test')

labels = []
texts = []

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in os.listdir(dir_name):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)

sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
```



```
[28]: # evaluate the model on test data set
model.evaluate(x_test, y_test)
```

```
782/782 [=====] - 2s 2ms/step - loss: 0.7313 - acc:
0.5232
```

```
[28]: [0.7313319444656372, 0.5231599807739258]
```

Assginment 10.3

Using the LSTM layer in Keras Using listing 6.27 in Deep Learning with Python as a guide, fit the same data with an LSTM layer. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

```
[30]: # set variable
max_features = 10000
```

```
[31]: model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
```

```
[32]: model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
```

```
[33]: history = model.fit(x_train, y_train, epochs=10, batch_size=128,
    ↪ validation_data=(x_val, y_val))
```

Epoch 1/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6933 - acc: 0.4650 -
val_loss: 0.6927 - val_acc: 0.5179
```

Epoch 2/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6878 - acc: 0.7150 -
val_loss: 0.6920 - val_acc: 0.5338
```

Epoch 3/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6801 - acc: 0.8550 -
val_loss: 0.6909 - val_acc: 0.5663
```

Epoch 4/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6696 - acc: 0.9350 -
val_loss: 0.6890 - val_acc: 0.5732
```

Epoch 5/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6512 - acc: 0.9650 -
val_loss: 0.6840 - val_acc: 0.5756
```

Epoch 6/10

```
2/2 [=====] - 3s 1s/step - loss: 0.6051 - acc: 0.9750 -
val_loss: 0.6552 - val_acc: 0.6181
```

Epoch 7/10

```
2/2 [=====] - 2s 1s/step - loss: 0.5260 - acc: 0.8250 -
```

```
val_loss: 0.6571 - val_acc: 0.6022
Epoch 8/10
2/2 [=====] - 3s 1s/step - loss: 0.4027 - acc: 0.9300 -
val_loss: 0.6416 - val_acc: 0.6278
Epoch 9/10
2/2 [=====] - 3s 1s/step - loss: 0.2983 - acc: 0.9650 -
val_loss: 0.6588 - val_acc: 0.6393
Epoch 10/10
2/2 [=====] - 3s 1s/step - loss: 0.2593 - acc: 0.9600 -
val_loss: 0.6788 - val_acc: 0.6385
```

Plot the data



```
[34]: acc = history.history['acc']
      val_acc = history.history['val_acc']
      loss = history.history['loss']
      val_loss = history.history['val_loss']
      epochs = range(1, len(acc) + 1)

      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and Validation of Accuracy')
      plt.legend()

      plt.figure()

      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and Validation of Loss')
      plt.legend()

      plt.show()
```



Evaluate the model on test set

```
[35]: model.evaluate(x_test, y_test)
```

```
782/782 [=====] - 15s 19ms/step - loss: 0.6837 - acc: 0.6322
```

```
[35]: [0.6837488412857056, 0.6321600079536438]
```

0.0.2 Assignment 10.4

Training and evaluating a simple 1D convnet on the IMDB data Using listing 6.46 in Deep Learning with Python as a guide, fit the same data with a simple 1D convnet. Produce the model performance metrics and training and validation accuracy curves within the Jupyter notebook.

```
[37]: model = Sequential()
model.add(layers.Embedding(max_features, 128, input_length=maxlen))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 100, 128)	1280000
conv1d (Conv1D)	(None, 94, 32)	28704
max_pooling1d (MaxPooling1D)	(None, 18, 32)	0
conv1d_1 (Conv1D)	(None, 12, 32)	7200
global_max_pooling1d (Global	(None, 32)	0
dense_3 (Dense)	(None, 1)	33

Total params: 1,315,937
Trainable params: 1,315,937
Non-trainable params: 0

```
[38]: model.
      ↪ compile(optimizer=RMSprop(lr=1e-4), loss='binary_crossentropy', metrics=['acc'])
```

```
[39]: history = model.fit(x_train, y_train, epochs=10, batch_size=128,
    ↪validation_data=(x_val, y_val))
```

```
Epoch 1/10
2/2 [=====] - 1s 299ms/step - loss: 1.1820 - acc:
0.5000 - val_loss: 1.0516 - val_acc: 0.5024
Epoch 2/10
2/2 [=====] - 1s 273ms/step - loss: 0.9587 - acc:
0.5000 - val_loss: 0.9895 - val_acc: 0.5024
Epoch 3/10
2/2 [=====] - 1s 250ms/step - loss: 0.8808 - acc:
0.5000 - val_loss: 0.9480 - val_acc: 0.5024
Epoch 4/10
2/2 [=====] - 0s 249ms/step - loss: 0.8273 - acc:
0.5000 - val_loss: 0.9174 - val_acc: 0.5024
Epoch 5/10
2/2 [=====] - 0s 238ms/step - loss: 0.7866 - acc:
0.5000 - val_loss: 0.8922 - val_acc: 0.5024
Epoch 6/10
2/2 [=====] - 0s 247ms/step - loss: 0.7529 - acc:
0.5000 - val_loss: 0.8697 - val_acc: 0.5024
Epoch 7/10
2/2 [=====] - 0s 229ms/step - loss: 0.7235 - acc:
0.5000 - val_loss: 0.8523 - val_acc: 0.5024
Epoch 8/10
2/2 [=====] - 1s 250ms/step - loss: 0.6990 - acc:
0.5000 - val_loss: 0.8375 - val_acc: 0.5024
Epoch 9/10
2/2 [=====] - 0s 227ms/step - loss: 0.6765 - acc:
0.5000 - val_loss: 0.8220 - val_acc: 0.5024
Epoch 10/10
2/2 [=====] - 1s 267ms/step - loss: 0.6556 - acc:
0.5000 - val_loss: 0.8101 - val_acc: 0.5024
```

Plotting the data


```
[40]: acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)


plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and Validation of accuracy')
plt.legend()

plt.figure()
```

```
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and Validation of Loss')
plt.legend()

plt.show()
```





Evaluate the test data

```
[41]: model.evaluate(x_test, y_test)
```

```
782/782 [=====] - 3s 3ms/step - loss: 0.8107 - acc: 0.5000
```

```
[41]: [0.8106999397277832, 0.5]
```

```
[ ]:
```